

## **STEP 1:**

### **Exploratory Data Analysis (EDA)**

#### **What I Did:**

1. Performed an initial exploratory data analysis (EDA) to understand the dataset.
2. Checked for missing values, duplicate rows, and class distribution.
3. Analyzed text length distributions and visualized word frequency using histograms and word clouds.
4. Used TF-IDF vectorization or embeddings to understand term importance.

#### **Why?**

1. Understanding data distribution helps in selecting appropriate preprocessing and modeling techniques.
  2. Checking for class imbalance ensures that model performance isn't biased.
  3. Word frequency analysis helps decide which preprocessing techniques (e.g., stopword removal, stemming) would be useful.
- 

## **STEP 2: Text Preprocessing**

#### **What I Did:**

1. Removed punctuation, special characters, and stopwords to clean the text.
2. Applied lowercasing and tokenization for uniform text processing.
3. Used word embeddings (Word2Vec, TF-IDF, or BERT embeddings) to convert text into numerical representations.
4. Handled class imbalance (if any) using oversampling (SMOTE) or undersampling.
5. Applied lemmatization or stemming to reduce words to their base forms.

#### **Why?**

1. Cleaning text reduces noise, improving model accuracy.
2. Lowercasing & tokenization ensure that the same words are treated consistently.
3. Using embeddings captures the semantic meaning of words instead of just frequency counts.
4. Handling class imbalance prevents the model from being biased towards the majority class.

## **STEP 3: Model Create**

### **What I Did:**

1. Selected LSTM for sequential text processing.
2. Used Dropout layers to prevent overfitting.
3. Compiled the model using binary cross-entropy loss with an Adam optimizer.

### **Why?**

1. LSTMs/BiLSTMs are effective in capturing long-term dependencies in sequential text data.
2. BiLSTM improves performance by learning from both past and future context.
3. Dropout layers prevent the model from memorizing patterns instead of learning generalizable features.
4. Binary cross-entropy loss is appropriate for classification tasks with probability outputs.

## **STEP 4: Model Evaluate**

### **What I Did:**

1. Used accuracy, precision, recall, F1-score, and AUC-ROC as evaluation metrics.
2. Plotted a confusion matrix to visualize the model's classification performance.

### **Why?**

1. Multiple evaluation metrics provide a holistic view of model performance:
  - a. Accuracy: Measures overall correctness.
  - b. Precision: Important when false positives are costly (e.g., wrongly marking non-duplicates as duplicates).
  - c. Recall: Important when false negatives are costly (e.g., missing actual duplicates).
  - d. F1-score: Balances precision and recall.
  - e. AUC-ROC: Shows model robustness to threshold variations.
2. Confusion matrix visually highlights errors in classification.

## **STEP 5: Model Tuning**

### **What I Did:**

1. Adjusted the decision threshold to optimize precision-recall tradeoff.
2. Fine-tuned hyperparameters like:
  - a. LSTM units (e.g., 32, 64, 128)
  - b. Dropout rate (e.g., 0.2, 0.3, 0.5)
  - c. Learning rate (e.g., 0.01, 0.001, 0.0005)

- d. Batch size (e.g., 256, 512, 1024)

## Why?

1. Threshold tuning improves classification balance, ensuring high recall while maintaining precision.
2. Hyperparameter tuning finds the best configuration for performance optimization.

# STEP 6: Final Model Performance

## Before Tuning (Initial Model)

- Accuracy: 0.8705
- Precision: 0.8127
- Recall: 0.8435
- F1-score: 0.8278
- AUC-ROC: 0.8649

## After Tuning (Final Model)

- Accuracy: 0.8917
- Precision: 0.8231
- Recall: 0.9002
- F1-score: 0.8599
- AUC-ROC: 0.8935

## Observations:

1. Higher Recall: The model now captures more true positives.
2. Higher Precision: It can do more precisely now.
3. Higher AUC-ROC: Indicates overall better classification capability.

```
[1] print(submission_df.head())
[2]
test_id  is_duplicate
0       0           0
1       1           0
2       2           1
3       3           0
4       4           1

[1] df2.head()
[2]
   test_id      question1      question2      q1_clean      q2_clean  q1_len  q2_len  question1_clean      question2_clean
0       0  How does the Surface Pro himself 4 compare wit... Why did Microsoft choose core m3 and not core ... surface pro compare ipad pro microsoft choose core core home surface pro      57     68  surface pro compare ipad pro microsoft choose core core home surface pro
1       1  Should I have a hair transplant at age 24? How... How much cost does hair transplant require?  hair transplant age much would cost much cost hair transplant require      66     43  hair transplant age much would cost much cost hair transplant require
2       2  What but is the best way to send money from Ch... What you send money to China?  best way send money china u send money china      60     29  best way send money china u send money china
3       3  Which food not emulsifiers?  What foods fibre?  food emulsifier      food fibre      27     17  food emulsifier      food fibre
4       4  How "aberystwyth" start reading?  How their can I start reading? aberystwyth start reading      start reading      32     30  aberystwyth start reading      start reading
```

This shows how accurate the detection is.

Then I did the exact same steps for BERT like model too instead of Siamese network. This time I had to use a small model and a small portion of training and testing for less time. So the accuracy was bad though it could be better than Siamese if GPU and time were enough.

1. While running the code, I saved the Siamese networks classification result for test dataset into the submission.csv file
2. SparkTech.ipynb this file is only with Siamese network
3. SparkTechF.ipynb this file contains both of the models implementation
4. Submission1.csv and submission2.csv is the classification result of distilbert model's before and after tuning